



## Project Title: Real-Time Financial Fraud Detection with Explainable AI Dashboard

---

**Objective:** Build a real-time fraud detection system that:

- Ingests streaming transaction data
  - Predicts fraudulent activities with high recall
  - Explains model predictions using SHAP
  - Enables user feedback and model retraining
  - Deploys on a scalable, containerized infrastructure
- 

### Dataset:

- **Name:** PaySim (Synthetic Financial Transactions)
  - **Source:** Kaggle
  - **Reason:** Realistic structure, labeled, temporal and transactional behavior present
- 

### System Architecture Overview:

```
Frontend (React + shadcn)
  |
FastAPI Backend (REST APIs)
  |
Real-Time Scoring (Kafka/Redis Consumer + ML Model)
  |
Kafka/Redis Stream (Simulated Transaction Feed)
  |
Data Producer (pandas to stream)
```

---

### Tech Stack:

- **ML Models:** XGBoost, Autoencoder, GNN (optional), Temporal Transformer (optional)
  - **Backend:** FastAPI, SHAP, Redis/Kafka
  - **Frontend:** React.js, Tailwind CSS, shadcn/ui
  - **Streaming:** Apache Kafka / Redis Streams
  - **Explainability:** SHAP
  - **Containerization:** Docker
  - **Orchestration:** Kubernetes (Minikube -> AWS EKS)
  - **MLOps:** MLflow, Evidently AI, GitHub Actions
-

## Modeling & Evaluation:

- **XGBoost:** Tabular feature-based baseline
- **Autoencoder:** Unsupervised anomaly detection
- **Graph Neural Network:** User-to-user or user-to-merchant graph detection
- **Transformer:** Temporal sequential fraud patterns

## Evaluation Metrics:

- Recall (critical)
  - Precision, F1 Score
  - PR-AUC (for imbalanced classes)
- 

## Feature Engineering:

- Delta features: balance\_before - amount, etc.
  - Time features: hour of day, recency
  - Transaction type encoding
  - Graph features (if GNN used)
  - User behavior stats: avg. transaction size, frequency
- 

## Backend API:

- `POST /predict` : Returns fraud score + SHAP
  - `POST /feedback` : Accepts user-labeled data
  - `GET /stats` : Model + transaction summary metrics
- 

## Frontend Dashboard:

- Built with React + shadcn
  - **Components:**
    - Live Transaction Feed (color-coded fraud risk)
    - Explainability Modal (SHAP plot)
    - Feedback Buttons (False Positive/Negative)
    - Admin Analytics (charts and summaries)
    - Optional D3-based network graph for fraud rings
- 

## Deployment:

- **Docker:** Backend, frontend, inference service
- **Kubernetes:** Local (Minikube), cloud (EKS)
- **CI/CD:** GitHub Actions (build, test, deploy)
- **Monitoring:**
  - MLflow for experiments
  - Evidently AI for data drift
  - Prometheus + Grafana (optional)

---

### Folder Structure:

```
fraud-detection-project/  
├─ backend/  
├─ frontend/  
├─ streaming/  
├─ ml_pipeline/  
│   ├── eda.py  
│   ├── model_train.py  
│   └── shap_explain.py  
├─ docker/  
├─ k8s/  
└─ README.md
```

---

### Learning Outcomes:

- End-to-end ML pipeline development
- Real-time data engineering with Kafka
- SHAP-based explainability in production
- Full-stack ML system with FastAPI + React
- MLOps tools: MLflow, Docker, K8s, GitHub Actions
- Fraud-specific domain modeling and mitigation strategies

---

### Next Steps:

1. Complete EDA ( `eda.py` )
  2. Perform feature engineering and model training
  3. Build model inference and SHAP explanation API
  4. Connect Kafka/Redis with FastAPI pipeline
  5. Create React frontend with shadcn components
  6. Containerize and orchestrate
  7. Deploy with monitoring and feedback loop
-